

CHRIS McNOSKY

AI EVALUATION & RED-TEAMING · AI QUALITY & OPERATIONS ANALYSIS · TECHNOLOGY, VENDOR & COMPLIANCE RISK

Dallas-Fort Worth, TX · [linkedin.com/in/chris-mcnosky](https://www.linkedin.com/in/chris-mcnosky) · github.com/cmcnosky · cmcnosky.github.io/why-hire

SUMMARY

Adversarial AI evaluator and evidence-driven technical reviewer with a background in procurement and regulated operations. Separates what a system demonstrably does from what it is claimed to do, and converts the difference into explicit requirements, deterministic tests, and defensible release decisions. Brings law-school evidence training to the question an evaluation turns on: does the record support the claim, or only appear to. Directs AI agents under written controls, and retains problem definition, rubrics, evidence review, and release judgment.

SELECTED PROJECTS

Stinger — coding-agent integrity evaluator | github.com/cmcnosky/stinger | public repo · public CI on every push

- Designed 30 public validated scenarios across five violation families — test integrity, spec fidelity, claims audit, governance compliance, and safety/exfiltration — scored by seven deterministic detectors rather than an AI judge.
- Shipped 8 reproducible evidence packages, each pinning rubric version, corpus hash, and config fingerprint, and rerunning offline to identical per-scenario outcomes.
- Corrected a defect in the evaluator itself: when the harness mislabeled a correct refusal, preserved the original incorrect package unedited rather than restating the record, broadened the refusal markers, regression-tested against six independently worded refusals, and documented why the defect moved a behavior label and not the score.
- Held the work at benchmark-*candidate* status under an explicit release HOLD rather than claim a finished benchmark the evidence does not yet support.

Side Effects Lab — agent reliability under missing or contradictory tool evidence | in development

- Rejected a candidate change carrying 1,336 passing tests after exact-commit adversarial review found two false-pass classes where the verification oracle still returned PASS. Verdict: CHANGES_REQUESTED; merge held.

WASP 2.0 — clean-room Rust, Python, and PostgreSQL transaction-integrity prototype | directed; agents implemented

- Specified deterministic client-order identity, durable recovery from ambiguous submissions, append-only evidence schemas, reconcile-first startup, and human activation permits; missing or contradictory evidence fails closed rather than resolving in the system's favor.

EXPERIENCE

Founder / AI Evaluation & Governance — NoFuckery AI | 2026 – present

- Own the technical decisions end to end: define the problem and the architecture, set acceptance criteria and decision rights, review the evidence, and issue the release call — GO, HOLD, CORRECT, or STOP.
- Direct multi-agent engineering work under written controls and separation of duties, so the party that produced the work never certifies it; verify completion claims against commands, files, diffs, tests, and held-out evidence rather than a system's own report that it finished.
- Hold releases when the evidence does not support the claim, and keep the correction record — including the original finding — rather than restating it.

Founder / Sole Operator — High Pie | October 2022 – present

- Run a regulated hemp-product business end to end with no support staff, owning every call: sourcing and supplier selection, pricing, inventory, compliance documentation, fulfillment, and customer problems — in a category where the paperwork has to survive inspection.

Purchasing Coordinator — Lennar | May 2018 – October 2022

- Evaluated bids, contracts, budgets, and vendor records for a national homebuilder; reconciled cost and documentation discrepancies before they reached the field.
- Coordinated vendors, trades, and internal stakeholders against hard construction deadlines, balancing cost, quality, schedule, documentation, and compliance.
- Maintained accurate, deadline-sensitive purchasing documentation across a high volume of concurrent work.

SKILLS

Evaluation & red-teaming — adversarial scenario and prompt design; rubric and scorecard design with criteria frozen before execution; capability- and completion-claim verification; refusal and failure classification; transcript, output, log, and diff review.

Evidence & governance — reproducible evidence packages and failure reports; fail-closed control design; specification and acceptance-criteria drafting; release judgment (GO / HOLD / CORRECT / STOP); multi-agent orchestration with separation of duties; technical research and writing from primary sources — observation vs. reported claim vs. inference.

Operations & tooling — contract, bid, and vendor-record administration; procurement and third-party risk; regulated-product compliance documentation; technical translation for technical and nontechnical audiences; Git, GitHub, and CI; Python (intermediate); Excel (PivotTables, XLOOKUP, data validation, macros/VBA, Power Query).

EDUCATION

University of North Texas — B.A., Sociology, Economics, and History

Brooklyn Law School — Coursework toward J.D., 2014-2017; degree not conferred