

CHRIS McNOSKY

AI EVALUATION & QUALITY OPERATIONS

Dallas-Fort Worth, TX | 817-360-0687 | cmcnosky@gmail.com

linkedin.com/in/chris-mcnosky | github.com/cmcnosky | cmcnosky.github.io/why-hire

PROFESSIONAL SUMMARY

I test claims against evidence. I turn hidden assumptions, contradictory records, and false passes into explicit requirements, reproducible checks, correction paths, and clear GO / CORRECT / HOLD / STOP decisions. My background combines AI-agent evaluation and agentic engineering direction with earlier work in purchasing, vendor documentation, and regulated operations.

SELECTED EVIDENCE

- **Tokio contribution: 83 executed upstream CI checks passed.** I scoped and directed an agentic Rust change for a select! macro issue that had outlasted two prior implementation attempts. The submitted, ready-for-review PR removes cfg-disabled branches before generated storage and scheduling, preserving dense branch indices and randomized fairness.
- **1,336 tests passed. The integration decision stayed HOLD.** In a private Side Effects Lab review, I directed a separate adversarial review of an agent-produced commit. It found two high-priority cases where pass/fail logic could incorrectly approve the behavior. I kept the candidate out of integration.
- **The agent refused correctly. Stinger labeled it incorrectly.** I preserved the original evidence package, including the wrong label; specified the correct classification; and required regression coverage for six independently worded refusals before accepting the correction.

PROFESSIONAL EXPERIENCE

Founder / AI Evaluation & Governance - NoFuckery AI | 2026-Present

- Define evaluation objectives, architecture constraints, written acceptance criteria, and decision rights. Direct AI-agent implementation; inspect files, diffs, logs, tests, reports, and source material; issue final GO / CORRECT / HOLD / STOP decisions.
- Separate production from review. The agent that produces the work does not certify it. Reject unsupported completion claims and preserve contradictory evidence for correction.
- Direct Stinger, a fail-closed Python evaluator with 30 public development and conformance scenarios and seven deterministic detectors for AI-agent integrity.

Founder / Sole Operator - High Pie | October 2022-Present

- Own supplier selection, sourcing, production, pricing, inventory, regulated-product compliance documentation, fulfillment, and customer issue resolution for a hemp-product business.

Purchasing Coordinator - Lennar | May 2018-October 2022

- Reviewed bids, contracts, budgets, and vendor records; reconciled cost and documentation discrepancies before field execution.
- Coordinated vendors, trades, and internal stakeholders against construction deadlines across cost, quality, schedule, documentation, and compliance.

CORE SKILLS

AI evaluation and quality operations: adversarial scenario design; acceptance criteria; rubric and scorecard design; completion-claim verification; transcript, output, log, and diff review; refusal and failure classification; correction and retest gates; agent red teaming

Evidence and governance: evidence review; separation of duties; fail-closed controls; release gates; primary-source technical research and writing

Operations and tools: bids, contracts, budgets, and vendor records; discrepancy resolution; regulated-product documentation; Python; Git, GitHub, and CI; Excel PivotTables, XLOOKUP, data validation, VBA, Power Query, and Power Pivot

EDUCATION

University of North Texas - B.A., Sociology, Economics, and History

Brooklyn Law School - Coursework toward J.D., 2014-2017; degree not conferred